

# LING 571 Final Project

---

Author: Dylan Watson

Professor: Rob Malouf

Course: LING 571: Computational Linguistics

## Introduction:

For this project, I will be using methods of computational linguistics and machine learning to categorize news headlines and text as well as analyzing linguistic differences between categories.

## Question:

How can we classify news text into categories using computational linguistic techniques and what are the linguistic differences between these categories?

## The Dataset:

The dataset I have chosen contains data on 210k news articles written by HuffPost between the years 2012 and 2022. The dataset is provided by a user named Rishabh Misra on Kaggle and will be accredited in the citations at the bottom.

## Import Libraries:

```
In [25]: # Had to pip everything b/c opened Jupyter in new Anacondas environment running Python 3
# because tensorflow version >2.10 doesn't support gpu acceleration on Windows 10... :(
# Needs CUDA 11.2.0 and cuDNN 8.1.0
#!pip install tensorflow==2.10.0
#!pip install tensorflow
#!pip install pandas
#!pip install matplotlib
#!pip install nltk
#!pip install scikit-learn
#!pip install gensim
#!pip install nbconvert
#!pip install pyppeteer
```

```
In [21]: import pandas as pd
import json
import matplotlib.pyplot as plt
import string
import numpy as np
from collections import Counter

import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer, PorterStemmer
```

```

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM, Embedding, Flatten, Dropout
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.callbacks import EarlyStopping

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.feature_extraction.text import TfidfVectorizer

from gensim.models import Word2Vec

```

```

In [3]: print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
        print(tf.__version__)

```

```

Num GPUs Available:  1
2.10.0

```

## Import Data:

The data in the json file is organized as individual json objects per article/row and must be read in row by row to a dataframe instead of all at once.

```

In [6]: data = []
        with open('News_Category_Dataset_v3.json', 'r') as file:
            for line in file:
                data.append(json.loads(line))
        news_data_df = pd.DataFrame(data)

```

## Exploratory Data Analysis:

Before we begin with any manipulation of the data, lets take a look at what data we have and how its structured.

```

In [5]: print("Head:")
        print("-----")
        print(news_data_df.head())
        print("\nDescribe:")
        print("-----")
        print(news_data_df.describe())
        print("\nInfo:")
        print("-----")
        print(news_data_df.info())
        print("\nData Types:")
        print("-----")
        print(news_data_df.dtypes)
        print("\nMissing:")
        print("-----")
        print(news_data_df.isnull().sum())
        print("\nColumns:")
        print("-----")
        print(news_data_df.columns)

```

```

Head:

```

```

-----

```

```

link \

```

```

0  https://www.huffpost.com/entry/covid-boosters-...

```

```

1  https://www.huffpost.com/entry/american-airlin...

```

```

2 https://www.huffpost.com/entry/funniest-tweets...
3 https://www.huffpost.com/entry/funniest-parent...
4 https://www.huffpost.com/entry/amy-cooper-lose...

```

	headline	category
0	Over 4 Million Americans Roll Up Sleeves For O...	U.S. NEWS
1	American Airlines Flyer Charged, Banned For Li...	U.S. NEWS
2	23 Of The Funniest Tweets About Cats And Dogs ...	COMEDY
3	The Funniest Tweets From Parents This Week (Se...	PARENTING
4	Woman Who Called Cops On Black Bird-Watcher Lo...	U.S. NEWS

	short_description	authors
0	Health experts said it is too early to predict...	Carla K. Johnson, AP
1	He was subdued by passengers and crew when he ...	Mary Papenfuss
2	"Until you have a dog you don't understand wha...	Elyse Wanshel
3	"Accidentally put grown-up toothpaste on my to...	Caroline Bologna
4	Amy Cooper accused investment firm Franklin Te...	Nina Golgowski

	date
0	2022-09-23
1	2022-09-23
2	2022-09-23
3	2022-09-23
4	2022-09-22

Describe:

-----

	link	headline
count	209527	209527
unique	209486	207996
top	https://www.huffingtonpost.comhttps://www.wash...	Sunday Roundup
freq	2	90

	category	short_description	authors	date
count	209527	209527	209527	209527
unique	42	187022	29169	3890
top	POLITICS			2014-03-25
freq	35602	19712	37418	100

Info:

-----

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209527 entries, 0 to 209526
Data columns (total 6 columns):

```

#	Column	Non-Null Count	Dtype
0	link	209527 non-null	object
1	headline	209527 non-null	object
2	category	209527 non-null	object
3	short_description	209527 non-null	object
4	authors	209527 non-null	object
5	date	209527 non-null	object

```

dtypes: object(6)
memory usage: 9.6+ MB
None

```

Data Types:

-----

link	object
headline	object
category	object
short_description	object
authors	object
date	object
dtype:	object

```

Missing:
-----
link                0
headline            0
category            0
short_description   0
authors            0
date               0
dtype: int64

Columns:
-----
Index(['link', 'headline', 'category', 'short_description', 'authors', 'date'], dtype='object')

```

## EDA Findings:

**Columns:** We find that we have 6 columns, they are link, headline, category, short\_description, authors, and date.

**Missing:** There are no missing values. This dataset is clean! :)

**Data Types:** All datatypes of these columns are objects. Some may need to be converted to strings for text analysis.

**Describe:** We find that there are exactly 209,527 articles/rows. There also may be 19,712 articles that lack a short\_description despite not appearing as null. Will have to check and possibly remove those.

**Head:** Finally, here we can see what the data of each column actually looks like for a more intuitive understanding of it.

## Target Feature (Category) Bar Chart:

For the classification of news articles according to a linguistic analysis of their text, we will use the "category" column as our target value in modeling as well as how we will separate the articles for other analysis.

```
In [6]: category_counts = news_data_df['category'].value_counts()
```

```
In [7]: print("Category Counts:")
print("-----")
print(category_counts)
print("\nNum of Categories")
print("-----")
print(len(category_counts))
```

```

Category Counts:
-----
category
POLITICS          35602
WELLNESS          17945
ENTERTAINMENT     17362
TRAVEL            9900
STYLE & BEAUTY    9814
PARENTING         8791
HEALTHY LIVING    6694
QUEER VOICES      6347
FOOD & DRINK      6340
BUSINESS          5992
COMEDY            5400

```

SPORTS	5077
BLACK VOICES	4583
HOME & LIVING	4320
PARENTS	3955
THE WORLDPOST	3664
WEDDINGS	3653
WOMEN	3572
CRIME	3562
IMPACT	3484
DIVORCE	3426
WORLD NEWS	3299
MEDIA	2944
WEIRD NEWS	2777
GREEN	2622
WORLDPOST	2579
RELIGION	2577
STYLE	2254
SCIENCE	2206
TECH	2104
TASTE	2096
MONEY	1756
ARTS	1509
ENVIRONMENT	1444
FIFTY	1401
GOOD NEWS	1398
U.S. NEWS	1377
ARTS & CULTURE	1339
COLLEGE	1144
LATINO VOICES	1130
CULTURE & ARTS	1074
EDUCATION	1014

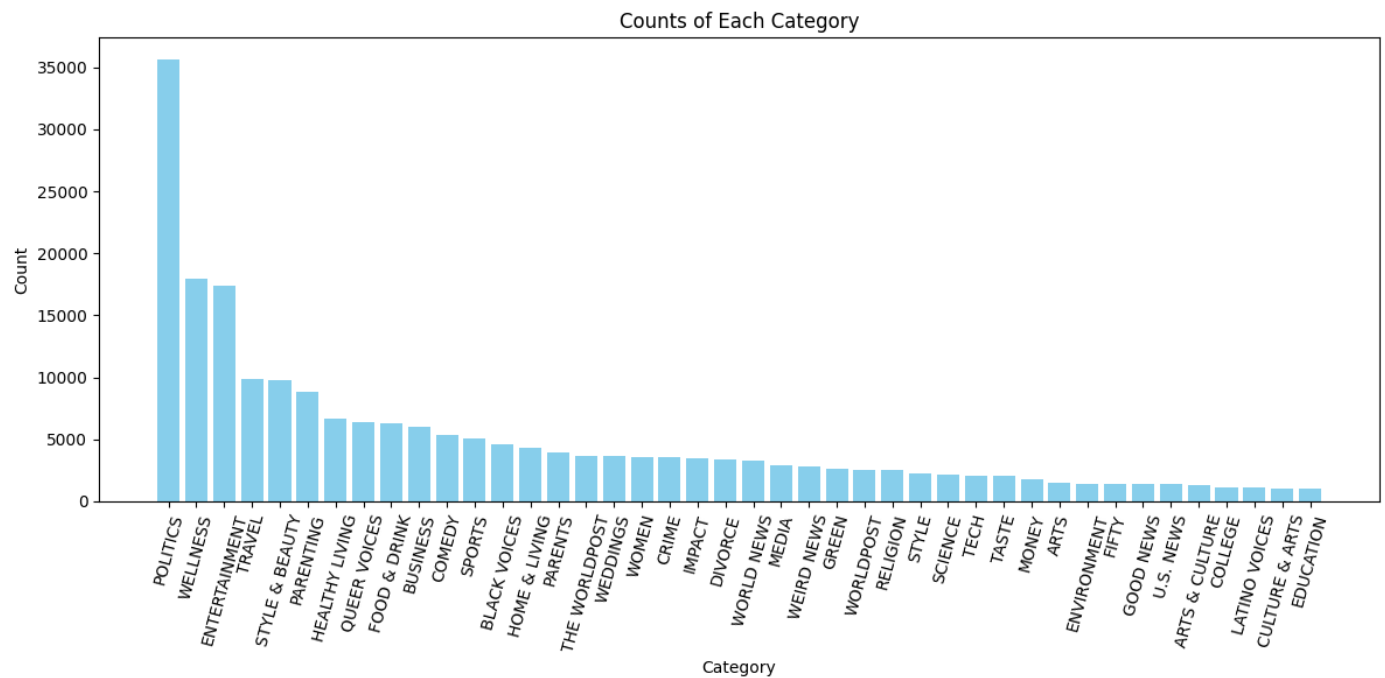
Name: count, dtype: int64

Num of Categories

-----

42

```
In [8]: categories = category_counts.index
counts = category_counts.values
plt.figure(figsize=(12, 6))
plt.bar(categories, counts, color='skyblue')
plt.title('Counts of Each Category')
plt.xlabel('Category')
plt.ylabel('Count')
plt.xticks(rotation=75)
plt.tight_layout()
plt.show()
```



## Target Feature (Category) Findings:

There are 42 unique categories. They do not seem to overlap or need any cleaning as they are consistent in format. The most popular in terms of frequency is POLITICS and the least is EDUCATION. The distribution of frequencies is right skewed and falls off rapidly approaching an asymptote of 10k. These are all the values we will be attempting to predict based on text analysis.

## Cleaning:

Lets check to see if some of those short\_descriptions are actually blank and remove them. Lets also remove the date, authors, and link features b/c the focus of this project is on computational linguistics and we are most interested in the headline, short\_description, and category features.

```
In [7]: columns_to_drop = ['link', 'authors', 'date']
news_data_cleaned_df = news_data_df.drop(columns=columns_to_drop)
```

```
In [8]: news_data_cleaned_df = news_data_cleaned_df[(news_data_cleaned_df['headline'] != '') & (
print(news_data_cleaned_df.describe())
```

	headline	category
count	189814	189814
unique	188417	42
top	Sunday Roundup	POLITICS
freq	90	32441

	short_description
count	189814
unique	187020
top	Welcome to the HuffPost Rise Morning Newsbrief...
freq	192

By removing all rows in which the headline or short\_description features are empty (but not null) we reduced the number of articles / rows by 19,713 to a new number of 189,814.

## Text Processing:

Here we need to tokenize, lowercase, remove punctuation, remove stopwords, stem, lemmatize, and normalize the text before we can feed it to a model.

Lowercasing - converting all characters to lowercase.

Remove Punctuation - as implied. Remove commas, periods, and all other symbols.

Remove Stopwords - stopwords are commonly used words in the english language that are mostly for grammatical purposes but don't contain as much value individually.

Stemming - convert words down to their base word. Plurals, words that end with 'ing', 'ed', 'es', 'er' May not be a word.

Lemmatize - group different inflections of words to same word. Still keeps some context though unlike stemming.

(yes I know stemming and lemmatizing do quite similar things to text but I really needed to reduce the amount of data as much as possible so my computer could train a model on it, my cpu is crying)

```
In [12]: #nltk.download('punkt')
#nltk.download('wordnet')
#nltk.download('stopwords')
```

```
In [13]: def preprocess_text(text):
# Tokenization
tokens = word_tokenize(text.lower())
# Remove stopwords and punctuation
stop_words = set(stopwords.words('english'))
punctuation = set(string.punctuation)
tokens = [token for token in tokens if token.isalpha() and token not in stop_words a
# Lemmatization
lemmatizer = WordNetLemmatizer()
tokens = [lemmatizer.lemmatize(token) for token in tokens]
# Stemming
stemmer = PorterStemmer()
tokens = [stemmer.stem(token) for token in tokens]
return ' '.join(tokens)
```

\*The cell below takes a minute!

```
In [14]: news_data_processed_df = news_data_cleaned_df.copy()
news_data_processed_df['headline'] = news_data_cleaned_df['headline'].apply(preprocess_t
news_data_processed_df['short_description'] = news_data_cleaned_df['short_description'].
```

## Comparing Processed & Unprocessed Text:

```
In [14]: print("Cleaned Original Text:")
print("-----")
print(news_data_cleaned_df.head())
print("\nProcessed Text:")
print("-----")
print(news_data_processed_df.head())
```

Cleaned Original Text:

-----

	headline	category	\
0	Over 4 Million Americans Roll Up Sleeves For O...	U.S. NEWS	
1	American Airlines Flyer Charged, Banned For Li...	U.S. NEWS	
2	23 Of The Funniest Tweets About Cats And Dogs ...	COMEDY	
3	The Funniest Tweets From Parents This Week (Se...	PARENTING	
4	Woman Who Called Cops On Black Bird-Watcher Lo...	U.S. NEWS	

```

                                short_description
0  Health experts said it is too early to predict...
1  He was subdued by passengers and crew when he ...
2  "Until you have a dog you don't understand wha...
3  "Accidentally put grown-up toothpaste on my to...
4  Amy Cooper accused investment firm Franklin Te...

Processed Text:
-----

                                headline      category \
0          million american roll sleeve covid booster  U.S. NEWS
1  american airlin flyer charg ban life punch fli...  U.S. NEWS
2                                funniest tweet cat dog week      COMEDY
3                                funniest tweet parent week  PARENTING
4                                woman call cop black lose lawsuit  U.S. NEWS

                                short_description
0  health expert said earli predict whether deman...
1  subdu passeng crew fled back aircraft confront...
2                                dog understand could eaten
3  accident put toothpast toddler toothbrush scre...
4  ami cooper accus invest firm franklin templeto...

```

## Text Processing Result:

After processing the text we can verify that the columns are still aligned and can see that all the appropriate text manipulation has worked. For example "23 Of The Funniest Tweets About Cats And Dogs" becomes simplified down to "funniest tweet cat dog week." This reduces the total number of unique tokens while still maintaining most of the valuable context allowing us to more efficiently train a model on the data. Lets see exactly how much this rtext processing has reduced the amount of unique tokens.

## Processed Text Analysis:

```

In [15]: news_data_cleaned_combined_df = pd.DataFrame(columns=['combined_text', 'category'])
news_data_cleaned_combined_df['category'] = news_data_cleaned_df['category'].copy()
news_data_cleaned_combined_df['combined_text'] = news_data_cleaned_df['headline'] + ' '

news_data_processed_combined_df = pd.DataFrame(columns=['combined_text', 'category'])
news_data_processed_combined_df['category'] = news_data_processed_df['category'].copy()
news_data_processed_combined_df['combined_text'] = news_data_processed_df['headline'] +

```

```

In [16]: cleaned_combined_all = news_data_cleaned_combined_df['combined_text'].str.cat(sep=' ')
processed_combined_all = news_data_processed_combined_df['combined_text'].str.cat(sep=' ')
print("Total Num of Cleaned Original Tokens")
print("-----")
print(len(cleaned_combined_all))
print("\nTotal Num of Unique Cleaned Original Tokens")
print("-----")
print(len(set(word_tokenize(cleaned_combined_all))))

print("\n\nTotal Num of Processed Tokens")
print("-----")
print(len(processed_combined_all))
print("\nTotal Num of Unique Processed Tokens")
print("-----")
print(len(set(word_tokenize(processed_combined_all))))

```

Total Num of Cleaned Original Tokens

-----

35455470



Total Num of Unique Cleaned Original Tokens

-----

155651

Total Num of Processed Tokens

-----

21113280

Total Num of Unique Processed Tokens

-----


57561

## Amount of Text Reduction:

The text processing reduced the number of word tokens by 40.5% from 35.5 million to 21.1 million. Even more important, the number of unique tokens was reduced by 63.0% from 155.7k down to 57.6k.

## Most Frequent Tokens:

```
In [17]: word_freq = Counter(word_tokenize(processed_combined_all))
top_n = 40
top_words = dict(word_freq.most_common(top_n))
plt.figure(figsize=(10, 6))
plt.bar(top_words.keys(), top_words.values(), color='skyblue')
plt.xlabel('Words')
plt.ylabel('Frequency')
plt.title('Top ' + str(top_n) + ' Most Frequent Tokens (After Cleaning & Processing)')
plt.xticks(rotation=75)
plt.tight_layout()
plt.show()
```



## Token Findings:

In this chart we find the most frequent tokens after processing the text. As we are looking at news articles and with politics being the most frequent category, there is no surprise that the token 'Trump' lands first in

the list due to his highly controversial nature during his presidency.

## Splitting:

Lets split the data into two sets, a training set to train the model and a test set to evaluate the performance of the model. Additionally, we'll split the training and test sets into x and y variables to represent the predictive and target features.

```
In [19]: X_train, X_test, y_train, y_test = train_test_split(news_data_processed_combined_df['com
```

## Model Creation / Training:

For our model I will try to use a recurrent neural network using Tensorflow to output a confusion matrix of the probabilities of each category for each row of combined\_text. The category with the greatest probability will be assigned as the predicted category of the combined\_text. I've also implemented early stopping so that if the loss and validation loss begin to diverse due to overfitting that the training will stop. Additonally, I've plotted the loss and validation loss to see the training progress visually.

```
In [27]: # Tokenize
max_words = 10000
tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(X_train)

# Sequence
X_train_seq = tokenizer.texts_to_sequences(X_train)
X_test_seq = tokenizer.texts_to_sequences(X_test)

# Pad
max_len = 100
X_train_pad = pad_sequences(X_train_seq, maxlen=max_len)
X_test_pad = pad_sequences(X_test_seq, maxlen=max_len)

# Encode
label_encoder = LabelEncoder()
y_train_encoded = label_encoder.fit_transform(y_train)
y_test_encoded = label_encoder.transform(y_test)

# Number of classes
num_classes = 42

# RNN Architecture
model = Sequential()
model.add(Embedding(input_dim=max_words, output_dim=32, input_length=max_len))
# model.add(Flatten())
model.add(LSTM(64, return_sequences=False))
# model.add(Dense(64, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

# Compile
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accu

# Early Stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)

history = model.fit(X_train_pad, y_train_encoded, epochs=50, batch_size=64,
                    validation_data=(X_test_pad, y_test_encoded), callbacks=[early_stopp
```

```
# Plotting Loss and Validation Loss
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

Epoch 1/50

2373/2373 [=====] - 20s 8ms/step - loss: 2.3432 - accuracy: 0.4097 - val\_loss: 1.8013 - val\_accuracy: 0.5279

Epoch 2/50

2373/2373 [=====] - 20s 8ms/step - loss: 1.7115 - accuracy: 0.5570 - val\_loss: 1.6088 - val\_accuracy: 0.5702

Epoch 3/50

2373/2373 [=====] - 19s 8ms/step - loss: 1.5308 - accuracy: 0.5967 - val\_loss: 1.5264 - val\_accuracy: 0.5872

Epoch 4/50

2373/2373 [=====] - 19s 8ms/step - loss: 1.4226 - accuracy: 0.6203 - val\_loss: 1.4935 - val\_accuracy: 0.5967

Epoch 5/50

2373/2373 [=====] - 19s 8ms/step - loss: 1.3475 - accuracy: 0.6360 - val\_loss: 1.4878 - val\_accuracy: 0.5977

Epoch 6/50

2373/2373 [=====] - 20s 8ms/step - loss: 1.2896 - accuracy: 0.6488 - val\_loss: 1.4838 - val\_accuracy: 0.6008

Epoch 7/50

2373/2373 [=====] - 19s 8ms/step - loss: 1.2364 - accuracy: 0.6612 - val\_loss: 1.4782 - val\_accuracy: 0.6026

Epoch 8/50


2373/2373 [=====] - 19s 8ms/step - loss: 1.1922 - accuracy: 0.6710 - val\_loss: 1.5075 - val\_accuracy: 0.6023

Epoch 9/50

2373/2373 [=====] - 19s 8ms/step - loss: 1.1462 - accuracy: 0.6818 - val\_loss: 1.5249 - val\_accuracy: 0.6012

Epoch 10/50

2373/2373 [=====] - 19s 8ms/step - loss: 1.1070 - accuracy: 0.6915 - val\_loss: 1.5518 - val\_accuracy: 0.5983



## Evaluation:

Now that the model is trained, let's run the test dataset through the trained model and see how it performs.

```
In [28]: loss, accuracy = model.evaluate(X_test_pad, y_test_encoded)
y_pred_prob = model.predict(X_test_pad)
y_pred = np.argmax(y_pred_prob, axis=1)
conf_matrix = confusion_matrix(y_test_encoded, y_pred)
class_report = classification_report(y_test_encoded, y_pred)
print("Accuracy:", accuracy)
print("\nConfusion Matrix:\n", conf_matrix)
```

```
1187/1187 [=====] - 4s 3ms/step - loss: 1.4782 - accuracy: 0.60
26
1187/1187 [=====] - 3s 2ms/step
Accuracy: 0.602613091468811
```

```
Confusion Matrix:
[[ 63  3  1 ...  0  0  0]
 [ 11 44  5 ...  6  1  0]
 [  6  3 291 ...  2  3  0]
 ...
 [  0  2 13 ... 172  1  0]
 [  0  1  3 ...  1 203 12]
 [  4  1  1 ...  0 19 93]]
```

```
In [29]: y_test_original = label_encoder.inverse_transform(y_test_encoded)
y_pred_original = label_encoder.inverse_transform(y_pred)

# classification report
class_report = classification_report(y_test_original, y_pred_original)

# Print
print("\nClassification Report with Original Category Names:\n", class_report)
```

Classification Report with Original Category Names:					
	precision	recall	f1-score	support	
ARTS	0.34	0.36	0.35	173	
ARTS & CULTURE	0.35	0.17	0.23	259	
BLACK VOICES	0.50	0.34	0.40	862	
BUSINESS	0.46	0.44	0.45	947	
COLLEGE	0.33	0.31	0.32	178	
COMEDY	0.53	0.41	0.46	926	
CRIME	0.45	0.50	0.47	559	
CULTURE & ARTS	0.47	0.29	0.36	236	
DIVORCE	0.72	0.76	0.74	713	
EDUCATION	0.45	0.23	0.30	186	
ENTERTAINMENT	0.61	0.72	0.66	3009	
ENVIRONMENT	0.39	0.21	0.27	290	
FIFTY	0.36	0.04	0.07	230	
FOOD & DRINK	0.60	0.76	0.67	1258	
GOOD NEWS	0.24	0.06	0.09	187	
GREEN	0.34	0.45	0.39	395	
HEALTHY LIVING	0.48	0.15	0.23	1070	
HOME & LIVING	0.71	0.67	0.69	849	
IMPACT	0.32	0.27	0.29	593	
LATINO VOICES	0.61	0.29	0.39	198	
MEDIA	0.57	0.26	0.36	478	
MONEY	0.55	0.36	0.43	383	
PARENTING	0.51	0.72	0.60	1733	
PARENTS	0.57	0.19	0.28	695	
POLITICS	0.72	0.83	0.77	6464	
QUEER VOICES	0.72	0.69	0.70	1124	
RELIGION	0.55	0.39	0.45	345	
SCIENCE	0.48	0.42	0.45	359	
SPORTS	0.64	0.67	0.65	908	
STYLE	0.43	0.24	0.31	331	
STYLE & BEAUTY	0.77	0.81	0.79	2028	
TASTE	0.29	0.01	0.02	409	
TECH	0.50	0.44	0.47	431	
THE WORLDPOST	0.46	0.50	0.48	694	
TRAVEL	0.66	0.76	0.71	1918	
U.S. NEWS	0.17	0.01	0.02	274	
WEDDINGS	0.73	0.78	0.75	729	
WEIRD NEWS	0.32	0.37	0.34	483	
WELLNESS	0.59	0.79	0.67	3514	
WOMEN	0.43	0.26	0.33	659	
WORLD NEWS	0.39	0.32	0.35	632	
WORLDPOST	0.43	0.37	0.40	254	
accuracy			0.60	37963	
macro avg	0.49	0.42	0.43	37963	
weighted avg	0.58	0.60	0.58	37963	

## Classification Report Context:

Precision: the accuracy of positive predictions.

Recall: ratio of correctly predicted positives to actual positives.

F1-Score: mean of precision and recall.

Support: number of actual occurrences of category in dataset.

## Evaluation Findings:

This initial model was found to be 60.3% accurate. This model could be improved by vectorizing the text and or modifying the neural network architecture. Looking at the classification report we find that there are

some inequities in the model's ability to predict certain categories over others. In particular the model predicted some categories well such as "STYLE & BEAUTY" with an f1-score of 0.79, "WEDDINGS" at 0.75, "POLITICS" at 0.77, and "DIVORCE" at 0.74. On the other hand, the model struggled with categories like "U.S. NEWS" at 0.02, "GOOD NEWS" at 0.09, and "TASTE" at 0.02. Why is that? It's possible that the text processing ended up losing too much context preventing the model from effectively predicting some categories. Additionally, the frequency (support) of categories was not even. The categories that were higher in frequency (support) had more data for the model to learn while categories with less frequency provided an inadequate volume of data for the model to learn. This could potentially be solved by upscaling the categories so that the distribution of categories is uniform. However, from a linguistic perspective, some categories may have a more unique set of tokens to differentiate themselves from others. If the set of processed tokens of two categories is too similar it may be more challenging to differentiate them predictively. Lets look at what tokens that the model found to be most indicative for each category to learn more about each one.

## Categorical Linguistic Analysis:

Lets see what tokens were weighted highest for each category in the model to see if we can find trends. ... I spent hours trying to get this to work to no avail. I cannot seem to analyze weights in the model to determine what tokens correlate to what categories. Maybe if it were a simple linear model I could use the token coefficients to see their correlations to each category. Alternatively, I could have just looked at the already pre-labelled data to find words that stood out in each area with more time allotted.

```
In [37]: '''
embedding_weights = model.layers[0].get_weights()[0] # Assuming the embedding layer is
dense_layer_weights = model.layers[-1].get_weights()[0] # Weights of the Dense layer

word_index = tokenizer.word_index
reverse_word_index = {index: word for word, index in word_index.items()}

# Top weighted words for each category
for class_index in range(num_classes):
    class_name = label_encoder.inverse_transform([class_index])[0]
    print(f"Category {class_name}:")

    class_weights = embedding_weights[:, class_index]
    sorted_indices = np.argsort(class_weights)[::-1][:10]

    #top_words = [reverse_word_index[idx] for idx in sorted_indices]
    for idx in sorted_indices:
        if idx in reverse_word_index:
            top_words.append(reverse_word_index[idx])
    print(top_words)
    print("-----")

# TODO: fix the encoded categories to appropriate spots
'''
```

```
Out[37]: '\n# Get the weights from the embedding layer\nembedding_weights = model.layers[0].get_w
eights()[0] # Assuming the embedding layer is the first layer (index 0)\n\n# Get weight
s from the Flatten layer to the Dense layer\ndense_layer_weights = model.layers[-1].get_
weights()[0] # Weights of the Dense layer\n\nword_index = tokenizer.word_index\nreverse
_word_index = {index: word for word, index in word_index.items()}\n\n# Top weighted word
s for each category\nfor class_index in range(num_classes):\n    class_name = label_enco
der.inverse_transform([class_index])[0]\n    print(f"Category {class_name}:")\n    \n
class_weights = embedding_weights[:, class_index]\n    sorted_indices = np.argsort(class
_weights)[::-1][:10]\n    \n    #top_words = [reverse_word_index[idx] for idx in sorted_
indices]\n    for idx in sorted_indices:\n        if idx in reverse_word_index:\n
```

```
top_words.append(reverse_word_index[idx])\n    print(top_words)\n    print("-----\n-----")\n\n# TODO: fix the encoded categories to appropriate spots\n'
```

## Word Embedding using Word2Vec:

Lets use Word2Vec to vectorize the words for each category and the use PCA to reduce the dimensions to plot and visualize for each category.

```
In [63]: grouped = news_data_processed_combined_df.groupby('category')['combined_text'].apply(' '

# Train Word2Vec
word2vec_models = {}
for category, text in grouped[['category', 'combined_text']].values:
    word2vec_model = Word2Vec([text], vector_size=100, window=5, min_count=1, sg=0)
    word2vec_models[category] = word2vec_model

# Create Category Vectors
category_vectors = {}
for category, model in word2vec_models.items():
    words = [word for word in model.wv.index_to_key if word in model.wv]
    vectors = [model.wv[word] for word in words]
    category_vectors[category] = sum(vectors) / len(vectors)

# Apply PCA
pca = PCA(n_components=2)
category_vecs_pca = pca.fit_transform(list(category_vectors.values()))

# K-means Clustering
num_clusters = 5 # Adjust as needed
kmeans = KMeans(n_clusters=num_clusters, n_init = 10)
kmeans.fit(category_vecs_pca)


# Cluster Labels
cluster_labels = kmeans.labels_
grouped['cluster'] = cluster_labels
```

```
In [64]: colors = {0: 'red', 1: 'blue', 2: 'green', 3: 'orange', 4: 'purple', 5: 'yellow'}

# Plot
plt.figure(figsize=(12, 8))
for category, cluster, vec_pca in zip(grouped['category'], grouped['cluster'], category_
    plt.scatter(vec_pca[0], vec_pca[1], color=colors[cluster], label=category)
    plt.annotate(category, (vec_pca[0], vec_pca[1]), textcoords="offset points", xytext=

# Plot Labels
plt.xlabel('PCA Component 1')
plt.ylabel('PCA Component 2')
plt.title('Clustering of Categories based on Word Embeddings')
plt.show()
```

Clustering of Categories based on Word Embeddings



```
In [68]: # Apply PCA
pca = PCA(n_components=3)
category_vecs_pca = pca.fit_transform(list(category_vectors.values()))

# K-means Clustering
num_clusters = 5 # Adjust as needed
kmeans = KMeans(n_clusters=num_clusters, n_init = 10)
kmeans.fit(category_vecs_pca)

# Cluster Labels
cluster_labels = kmeans.labels_
grouped['cluster'] = cluster_labels
```

```
In [69]: colors = {0: 'red', 1: 'blue', 2: 'green', 3: 'orange', 4: 'purple', 5: 'yellow'} # Add


# Create Plot
fig = plt.figure(figsize=(12, 8))
ax = fig.add_subplot(111, projection='3d')

# Plot Categories
for category, cluster, vec_pca in zip(grouped['category'], grouped['cluster'], category_
    ax.scatter(vec_pca[0], vec_pca[1], vec_pca[2], color=colors[cluster], label=category)
    ax.text(vec_pca[0], vec_pca[1], vec_pca[2], category, color='black', fontsize=6)

# Plot Labels
ax.set_xlabel('PCA Component 1')
ax.set_ylabel('PCA Component 2')
ax.set_zlabel('PCA Component 3')
plt.title('Clustering of Categories based on Word Embeddings (3D)')
plt.show()
```



## Clustering of Categories based on Word Embeddings (3D)



```
In [75]: for index, row in grouped.iterrows():
          category = row['category']
          cluster = row['cluster']
          print(f"Category: {category}, Cluster: {cluster}")
```

```
Category: ARTS, Cluster: 3
Category: ARTS & CULTURE, Cluster: 1
Category: BLACK VOICES, Cluster: 1
Category: BUSINESS, Cluster: 1
Category: COLLEGE, Cluster: 2
Category: COMEDY, Cluster: 1
Category: CRIME, Cluster: 1
Category: CULTURE & ARTS, Cluster: 3
Category: DIVORCE, Cluster: 2
Category: EDUCATION, Cluster: 4
Category: ENTERTAINMENT, Cluster: 0
Category: ENVIRONMENT, Cluster: 4
Category: FIFTY, Cluster: 2
Category: FOOD & DRINK, Cluster: 3
Category: GOOD NEWS, Cluster: 2
Category: GREEN, Cluster: 3
Category: HEALTHY LIVING, Cluster: 3
Category: HOME & LIVING, Cluster: 0
```

Category: IMPACT, Cluster: 3  
 Category: LATINO VOICES, Cluster: 2  
 Category: MEDIA, Cluster: 1  
 Category: MONEY, Cluster: 2  
 Category: PARENTING, Cluster: 3  
 Category: PARENTS, Cluster: 2  
 Category: POLITICS, Cluster: 0  
 Category: QUEER VOICES, Cluster: 1  
 Category: RELIGION, Cluster: 1  
 Category: SCIENCE, Cluster: 4  
 Category: SPORTS, Cluster: 1  
 Category: STYLE, Cluster: 1  
 Category: STYLE & BEAUTY, Cluster: 3  
 Category: TASTE, Cluster: 1  
 Category: TECH, Cluster: 2  
 Category: THE WORLDPOST, Cluster: 3  
 Category: TRAVEL, Cluster: 0  
 Category: U.S. NEWS, Cluster: 4  
 Category: WEDDINGS, Cluster: 1  
 Category: WEIRD NEWS, Cluster: 4  
 Category: WELLNESS, Cluster: 0  
 Category: WOMEN, Cluster: 1  
 Category: WORLD NEWS, Cluster: 3  
 Category: WORLDPOST, Cluster: 3

## Word2Vec, PCA, KMeans Findings:

After using Word2Vec to vectorize the processed words for each of the categories, I used PCA to reduce the dimensions to 2 and 3 and used kmeans to group the categories according to the likeness of their word vectors. I found that k=5 to be an appropriate value based on the 2d plot. The points which appear closer and share a common cluster color are considered to be more like one another. I did both 2D and 3D PCA to see if the visualizations yielded different insights. Most of the cluster elements are the same between the two.

The categories fall into 5 clusters (from 3D PCA):

- 0: ENTERTAINMENT, HOME & LIVING, POLITICS, TRAVEL, WELLNESS, QUEER VOICES, RELIGION, TASTE
- 1: ARTS & CULTURE, BLACK VOICES, BUSINESS, COMEDY, MEDIA, SPORTS, STYLE, WEDDINGS, WOMEN
- 2: COLLEGE, DIVORCE, FIFTY, GOOD NEWS, LATINO VOICES, MONEY, PARENTS, TECH
- 3: ARTS, CULTURE & ARTS, FOOD & DRINK, HEALTHY LIVING, IMPACT, PARENTING, STYLE & BEAUTY, THE WORLDPOST, WORLD NEWS, WORLDPOST
- 4: EDUCATION, ENVIRONMENT, SCIENCE, U.S. NEWS, WEIRD NEWS

I found many points of interest among these graphs and cluster of categories. Some things I found predictable, and other less so. I found cluster 4 to be realistic as it grouped together EDUCATION, ENVIRONMENT, and SCIENCE which are all topics I would consider adjacent to one another. On the humorous side, I was enamored to see FIFTY (articles for people 50 and up), DIVORCE, and PARENTS together in cluster 2. These groupings indicate to me that in some sense, the vectorization and clustering was successful.

While some things may have fell into separate clusters due to the number of clusters chosen, I noticed that some similar topics were still close to one another by euclidian distance due to similarities. For example, ARTS and CULTURE & ARTS were in separate clustes in the 2D plot but were similar in their prinicple component values as the articles likely contained similar text.

Not all clusters and proximities made sense though. For example, the categories PARENTS and PARENTING

were both in separate clusters and far apart by their principal component values. Intuitively by their names, I would have assumed that those news article categories would have more overlap in text.

## Category Keywords using TF-IDF:

Lets try to analyze the key words that differentiate each category from one another. To do this we will use tf-idf (term frequency- inverse term frequency).

```
In [23]: grouped = news_data_processed_combined_df.groupby('category')['combined_text'].apply(' '

# TF-IDF Vectorization
tfidf = TfidfVectorizer()
tfidf_matrix = tfidf.fit_transform(grouped['combined_text'])

# Extracting keywords for each category
keywords_per_category = {}
for i, category in enumerate(grouped['category']):
    feature_names = tfidf.get_feature_names_out()
    feature_index = tfidf_matrix[i, :].nonzero()[1]
    tfidf_scores = zip(feature_index, [tfidf_matrix[i, x] for x in feature_index])
    sorted_tfidf_scores = sorted(tfidf_scores, key=lambda x: x[1], reverse=True)[:10] #
    keywords_per_category[category] = [feature_names[i] for i, _ in sorted_tfidf_scores]

# Displaying keywords for each category
for category, keywords in keywords_per_category.items():
    print(f"Category: {category}")
    print(f"Keywords: {' '.join(keywords)}")
    print()
```

Category: ARTS

Keywords: art, artist, new, nighter, one, work, year, music, theatr, opera

Category: ARTS & CULTURE

Keywords: art, artist, book, new, woman, trump, make, year, show, one

Category: BLACK VOICES

Keywords: black, polic, new, woman, say, life, peopl, said, year, white

Category: BUSINESS

Keywords: busi, compani, new, year, work, make, time, peopl, job, one

Category: COLLEGE

Keywords: colleg, student, univers, campu, educ, school, sexual, graduat, say, new

Category: COMEDY

Keywords: trump, donald, colbert, video, maher, show, jimmi, fallon, stephen, like

Category: CRIME

Keywords: polic, man, kill, shoot, say, suspect, said, offic, allegedli, arrest

Category: CULTURE & ARTS

Keywords: imageblog, art, photo, artist, new, work, exhibit, galleri, world, week

Category: DIVORCE

Keywords: divorc, marriag, ex, child, get, one, relationship, date, life, parent

Category: EDUCATION

Keywords: school, educ, student, teacher, child, colleg, charter, learn, devo, new

Category: ENTERTAINMENT

Keywords: new, show, star, film, say, trump, movi, one, year, get

Category: ENVIRONMENT

Keywords: anim, climat, week, photo, chang, video, new, world, weather, year

Category: FIFTY

Keywords: life, year, one, time, age, like, day, love, woman, make

Category: FOOD & DRINK

Keywords: recip, food, make, photo, best, day, cook, eat, one, like

Category: GOOD NEWS

Keywords: dog, get, help, man, life, day, one, make, love, rescu

Category: GREEN

Keywords: climat, chang, dog, new, year, peopl, one, world, water, anim

Category: HEALTHY LIVING

Keywords: life, health, way, peopl, one, make, time, get, new, thing

Category: HOME & LIVING

Keywords: home, photo, decor, idea, make, day, design, diy, craft, hous

Category: IMPACT

Keywords: peopl, world, day, woman, one, year, child, help, life, need

Category: LATINO VOICES

Keywords: latino, latina, immigr, trump, puerto, new, rico, say, one, peopl

Category: MEDIA

Keywords: trump, news, fox, medium, new, report, say, donald, time, cnn

Category: MONEY

Keywords: credit, tax, money, financi, year, get, time, card, bank, new

Category: PARENTING

Keywords: child, parent, kid, mom, babi, one, time, day, like, make

Category: PARENTS

Keywords: kid, parent, mom, child, day, babi, thing, time, one, dad

Category: POLITICS

Keywords: trump, donald, presid, gop, say, republican, democrat, senat, clinton, state

Category: QUEER VOICES

Keywords: gay, queer, lgbt, lgbtq, transgend, new, tran, peopl, lesbian, marriag

Category: RELIGION

Keywords: pope, christian, muslim, church, medit, peopl, franci, god, religi, spiritu

Category: SCIENCE

Keywords: nasa, new, scientist, space, may, studi, research, scienc, video, mar

Category: SPORTS

Keywords: nfl, game, olymp, team, player, win, footbal, nba, sport, first

Category: STYLE

Keywords: look, fashion, new, dress, style, beauti, like, make, hair, week

Category: STYLE & BEAUTY

Keywords: photo, style, fashion, look, week, pinterest, huffpoststyl, dress, stylelist, tumblr

Category: TASTE

Keywords: recip, make, food, delici, cook, eat, get, new, best, way

Category: TECH

Keywords: appl, new, facebook, iphon, googl, week, video, look, compani, get

Category: THE WORLDPOST

Keywords: attack, kill, say, syria, peopl, trump, isi, syrian, refuge, aleppo

Category: TRAVEL

Keywords: travel, photo, world, hotel, new, one, citi, best, get, day

Category: U.S. NEWS

Keywords: said, polic, new, say, shoot, peopl, state, kill, california, year

Category: WEDDINGS

Keywords: wed, bride, marriag, photo, coupl, marri, day, get, love, bridal

Category: WEIRD NEWS

Keywords: man, get, say, dog, one, polic, like, peopl, fark, go

Category: WELLNESS

Keywords: life, health, time, one, make, peopl, studi, way, get, new

Category: WOMEN

Keywords: woman, sexual, one, like, day, life, get, time, want, abort

Category: WORLD NEWS

Keywords: trump, said, korea, kill, say, attack, presid, north, peopl, new

Category: WORLDPOST

Keywords: world, year, countri, state, war, iran, one, israel, new, peopl

In the output of the cell above we are provided the keywords with the highest weights for each category from the tf-idf algorithm. At a quick glance, we can see how effectively that this method worked in differentiating important key words of each category.

Since we used stemming, many of the key words have been reduced down to sometimes challenging stems to identify which do not obviously represent recognizable english. Some of these words include 'fark'. Luckily, with some inference, most of the other stems can be made out to their original forms. It may be better to not use stemming and just use lemmatization for this type of analysis in the future.

While most categories contained predictable keywords, I thought one category in particular jumped out as humorous, no pun intended. The category 'COMEDY's top two keywords were 'trump' and 'donald.'

Some categories had predictable overlap. 'ARTS' and 'ARTS & CULTURE' both had four keywords in common among their top 10. They were as follows: art, artist, new, and year. I was impressed that the RNN created to predict categories earlier was capable of differentiating these two despite containing similar text keywords.

These keywords may provide insight into why the model struggled to predict some of the categories. The categories in which the model struggled most were GOOD NEWS, U.S. NEWS, and TASTE. We can now see that these categories either contained words common among all categories and or contained significant overlap with other particular categories. For example, TASTE shares many keywords with FOOD & DRINK. Likewise, U.S. NEWS and GOOD NEWS overlapped heavily with other WORLD and NEWS type categories.

Some keywords appeared among the top 10 of many categories. These words include: day, year, one, new, etc. It may be useful to make a list of the most frequent keywords across the categories and add that to a list of words to remove in the preprocessing to further enhance the differences between one another for predictive analysis.

## Citations:

1. Misra, Rishabh. "News Category Dataset." arXiv preprint arXiv:2209.11429 (2022).
2. Misra, Rishabh and Jigyasa Grover. "Sculpting Data for ML: The first act of Machine Learning." ISBN 9798585463570 (2021).